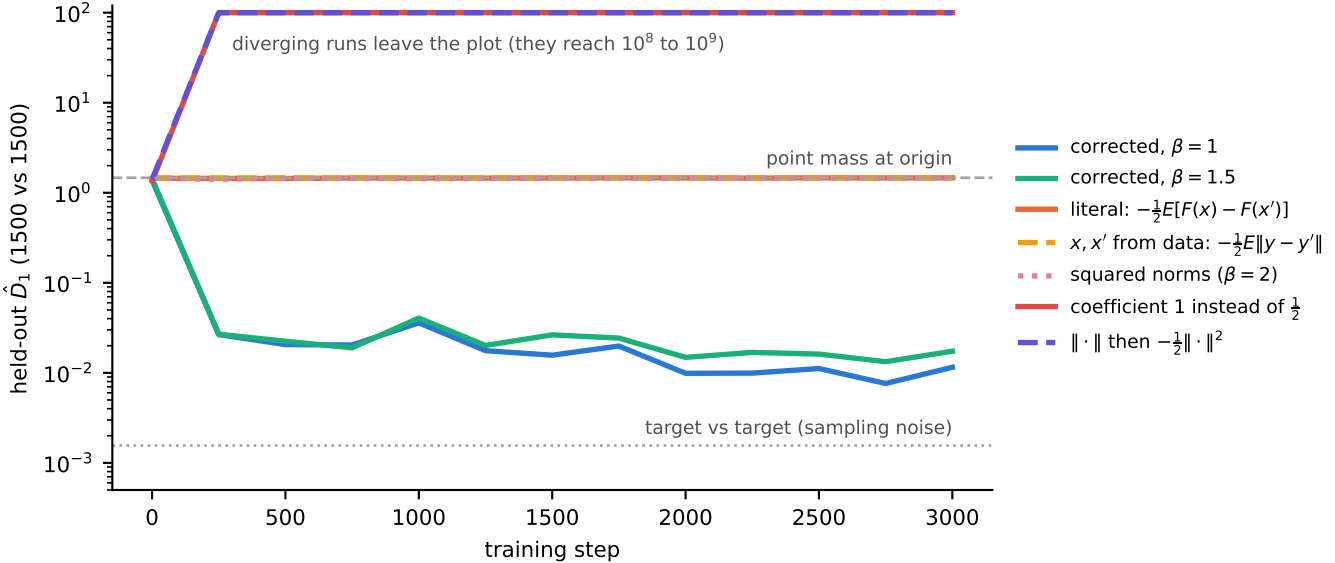
diverging runs leave the plot (they reach 10^8 to 10^9)
point mass at origin
target vs target (sampling noise)
held-out $\hat{D}_1$ (1500 vs 1500)
training step
10^2
10^1
10^0
10^{-1}
10^{-2}
10^{-3}
0
500
1000
1500
2000
2500
3000
corrected, $\beta = 1$
corrected, $\beta = 1.5$
literal: $-\frac{1}{2}E[F(x) - F(x')]$
x, x' from data: $-\frac{1}{2}E\|y - y'\|$
squared norms ($\beta = 2$)
coefficient 1 instead of $\frac{1}{2}$
$\|\cdot\|$ then $-\frac{1}{2}\|\cdot\|^2$